

UNIVERSITY OF CALIFORNIA, MERCED

# How to Get Your CVPR Paper Rejected?

Ming-Hsuan Yang

UCMERCED

[www.ucmerced.edu](http://www.ucmerced.edu)

# Outline

- Conferences
- Journals
- Writing
- Presentation
- Lessons

# Conferences

- CVPR – Computer Vision and Pattern Recognition, since 1983
  - Annual, held in US
- ICCV – International Conference on Computer Vision, since 1987
  - Every other year, alternate in 3 continents
- ECCV – European Conference on Computer Vision, since 1990
  - Every other year, held in Europe

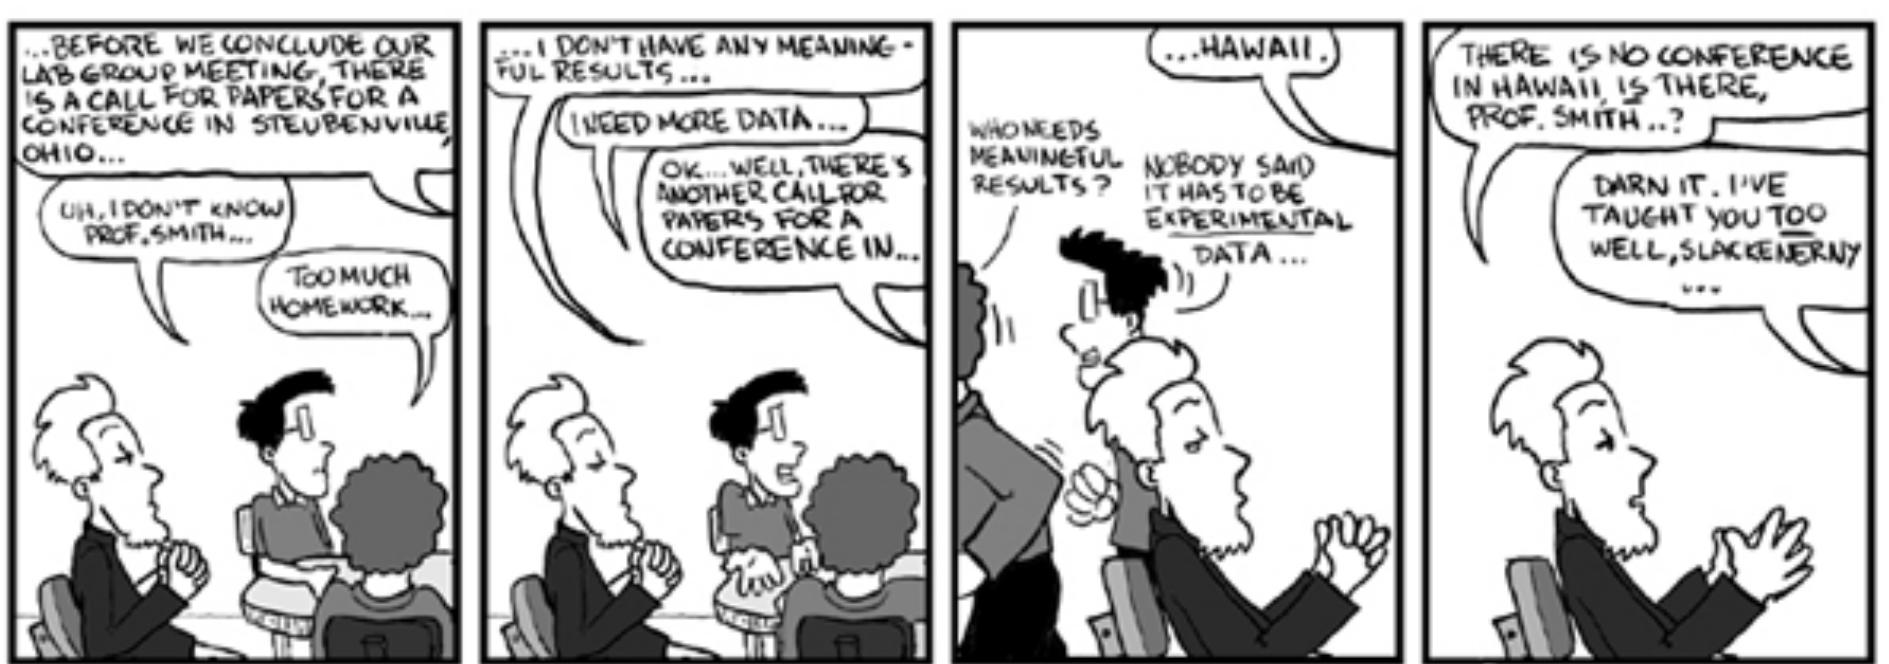
# Conferences

- ACCV – Asian Conference on Computer Vision
- BMVC – British Machine Vision Conference
- ICPR – International Conference on Pattern Recognition
- SIGGRAPH
- NIPS – Neural Information Processing Systems

# Conferences

- MICCAI – Medical Image Computing and Computer-Assisted Intervention
- FG – IEEE Conference on Automatic Face and Gesture Recognition
- ICCP – IEEE International Conference on Computational Photography
- ICML – International Conference on Machine Learning
- IJCAI, AAAI, MVA, ICDR, ICVS, DAGM, CAIP, ICRA, ICASSP, ICIP, SPIE, DCC, WACV, 3DPVT, ACM Multimedia, ICME, ...

# Conference Location



# Conference Location



- Me and conference I want to attend  
(location vs. reputation)

# Conference Organization

- General chairs: administration
- Program chairs: handling papers
- Area chairs:
  - Assign reviewers
  - Read reviews and rebuttals
  - Consolidation reports
  - Recommendation
- Reviewers
- Authors

# Review Process

- Submission
- CVPR/ECCV/ICCV
  - Double blind review
  - Program chairs: assign papers to area chairs
  - Area chairs: assign papers to reviewers
- Rebuttal
- Results

# Area Chair Meetings

- Each paper is reviewed by 2/3 area chairs
- Area chair make recommendations
- Program chairs make final decisions
- Virtual meetings
- Onsite meetings
  - Several panels
  - Buddy/triplet

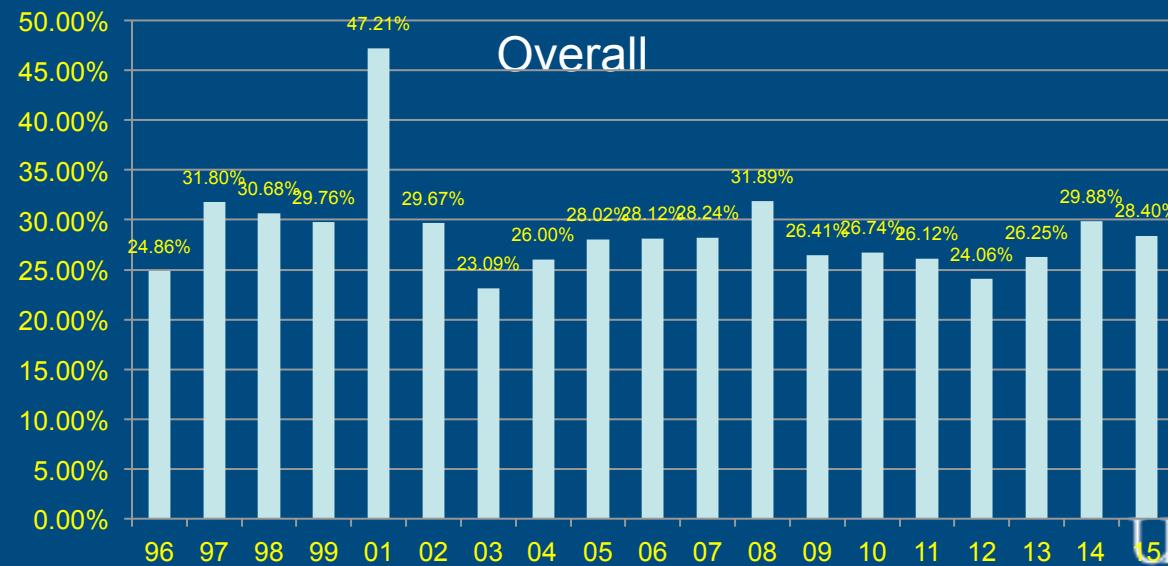
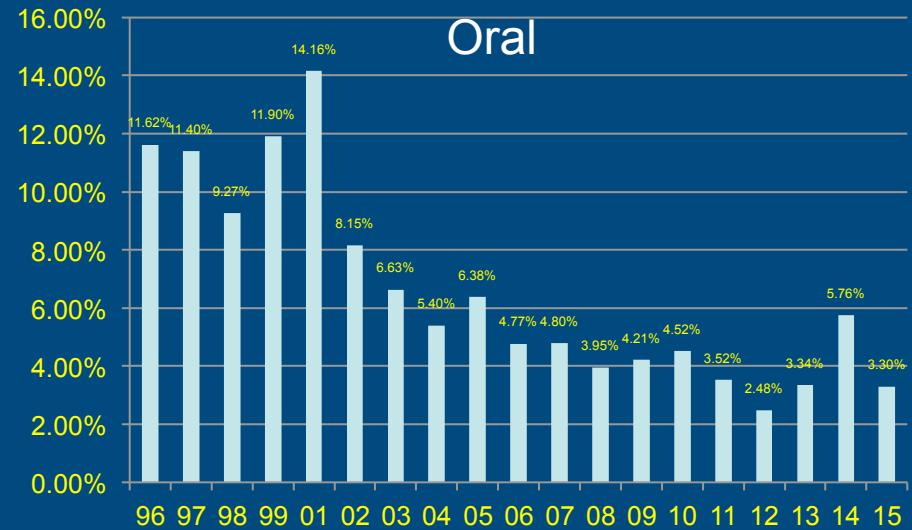
# Triage

- Area chairs know the reviewers
- Reviews are weighted
- Based on reviews and rebuttal
  - Accept: (decide oral later)
  - Reject: don't waste time
  - Go either way: lots of papers
- Usually agree with reviewers but anything can happen as long as there are good justifications

# Conference Acceptance Rate

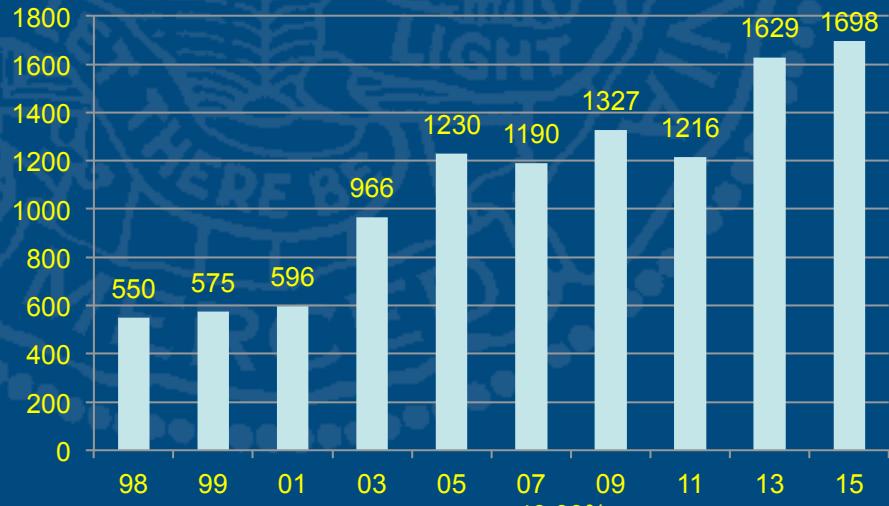
- ICCV/CVPR/ECCV: ~ 25%
- ACCV (2009): ~ 30%
- NIPS: ~ 25%
- BMVC: ~ 30%
- ICIP: ~ 45%
- ICPR: ~ 55%
- Disclaimer
  - low acceptance rate = high quality?

# CVPR

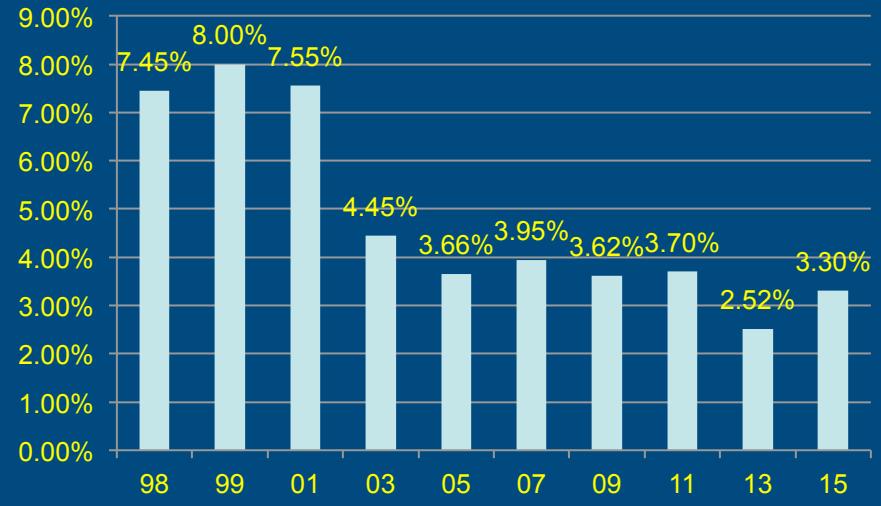


# ICCV

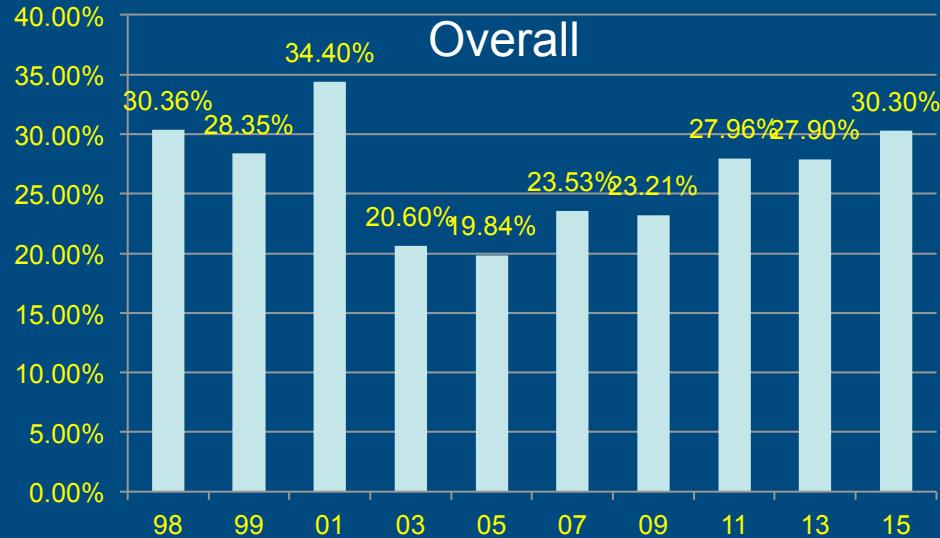
## Submission



## Oral

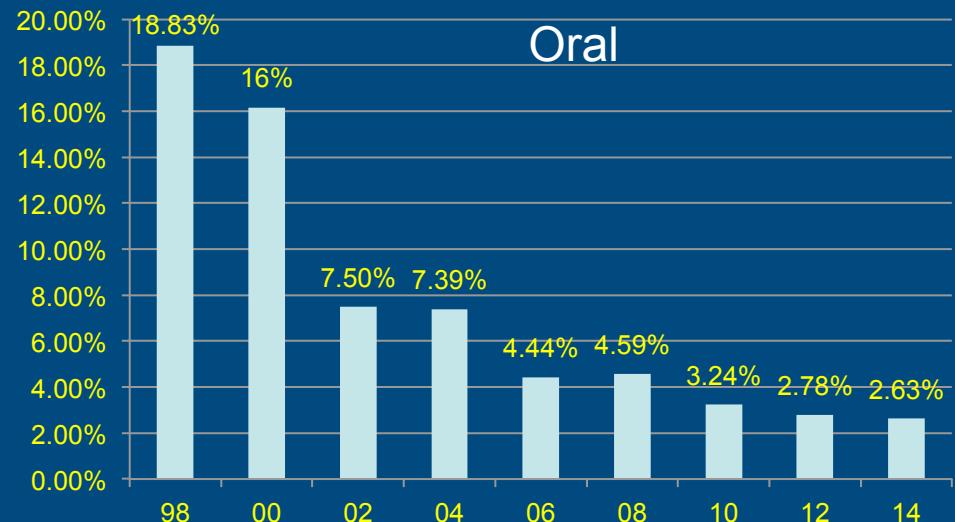
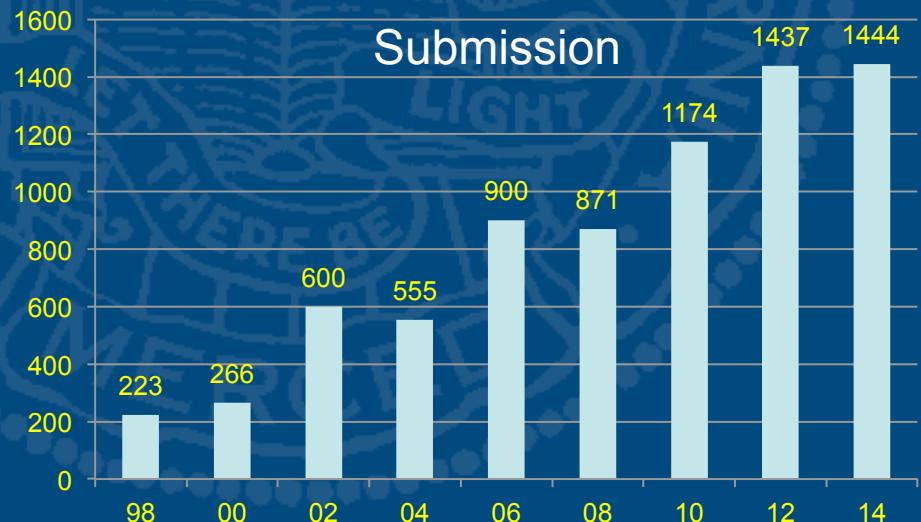


## Overall



UCMERCED

# ECCV



UCMERCED

# Top 100 Publications - English

- For what it is worth (h5 index by Google Scholar)
  1. Nature
  2. The New England Journal of Medicine
  3. Science
  - ...
  55. IEEE Conference on Computer Vision and Pattern Recognition (CVPR)

...

UCMERCED

# Top Publications - E&CS

1. Nano Letters

...

8. IEEE Conference on Computer Vision  
and Pattern Recognition (CVPR)

...

16. IEEE Transactions on Pattern Analysis  
and Machine Intelligence

...

# Reactions

- Top journal papers
- Workshops vs conferences
- Waiting for the review or final results
- Acceptance
- Reject
- Mixed feeling
- Finding an error
- Resubmit?
- This time, it will go through
- Paper finally accepted
- Registration
- Oral presentation
- Poster presentation

# Database Community

- Jeffrey Naughton's ICDE 2010 keynote
- What's wrong with the reviewing process?
- How to fix that?

# Journals

- PAMI – IEEE Transactions on Pattern Analysis and Machine Intelligence, since 1979 (impact factor: 5.96, #1 in all engineering and AI, top-ranked IEEE and CS journal)
- IJCV – International Journal on Computer Vision, since 1988 (impact factor: 5.36, #2 in all engineering and AI)
- CVIU – Computer Vision and Image Understanding, since 1972 (impact factor: 2.20)

# Journals

- IVC – Image and Vision Computing
- TIP – IEEE Transactions on Image Processing
- TMI- IEEE Transactions on Medical Imaging
- MVA – Machine Vision and Applications
- PR – Pattern Recognition
- TMM – IEEE Transactions on Multimedia
- ...

# PAMI Reviewing Process

- Associate editor-in-chief (AEIC) assigns papers to associate editors (AE)
- AE assigns reviewers
- First-round review: 2-4 months
  - Accept as is
  - Accept with minor revision
  - Major revision
  - Resubmit as new
  - Reject

# PAMI Reviewing Process

- Second-round review: 2-4 months
  - Accept as is
  - Accept with minor revision
  - Major revision (rare cases)
  - Reject
- EIC makes final decision
- Overall turn-around time: 6 to 12 months
- Rule of thumb: 30% additional work beyond a CVPR/ICCV/ECCV paper

# IJCV/CVIU Reviewing Process

- Similar formats
- Slightly longer turn-around time

# Journal Acceptance Rate

- PAMI
  - 2013: 151/959: 15.7%
  - 2014: 160/1018: 15.7%
- IJCV: ~ 20% (my guess, no stats)
- CVIU: ~ 25% (my guess, no stats)

# From Conferences to Journals

- How much additional work?
  - 30% additional more work for PAMI?
  - As long as the journal version is significantly different from the conference one
- Novelty of each work
  - Some reviewers still argue against this
  - Editors usually accept paper with the same ideas

# How to Get Your CVPR Paper Rejected?

- Jim Kajia (SIGGRAPH 93 papers chair):  
[How to get your SIGGRAPH paper rejected?](#)
- Bill Freeman:  
[How to write a good CVPR submission](#)
- Do not
  - Pay attention to review process
  - Put yourself as a reviewer to exam your work from that perspective
  - Put the work in right context
  - Carry out sufficient amount of experiments
  - Compare with state-of-the-art algorithms
  - Pay attention to writing

# Review Form

- Summary
- Overall Rating
  - Definite accept, weakly accept, borderline, weakly reject, definite reject
- Novelty
  - Very original, original, minor originality, has been done before
- Importance/relevance
  - Of broad interest, interesting to a subarea, interesting only to a small number of attendees, out of CVPR scope

# Review Form

- Clarity of presentation
  - Reads very well, is clear enough, difficult to read, unreadable
- Technical correctness
  - Definite correct, probably correct but did not check completely, contains rectifiable errors, has major problems
- Experimental validation
  - Excellent validation or N/A (a theoretical paper), limited but convincing, lacking in some aspects, insufficient validation
- Additional comments
- Reviewer's name

# Learn from Reviewing Process

- Learn how others/you can pick apart a paper
- Learn from other's mistakes
- Get to see other reviewers evaluate the same paper
- See how authors rebut comments
- Learn how to write good papers
- Learn what it takes to get a paper published

# Put Yourself as Reviewer

- Reviewer's perspective
- How a paper gets rejected?
- What are the contributions?
- Does it advance the science in the filed?
- Why you should accept this paper?
- Is this paper a case study?
- Is this paper interesting?
- Who is the audience?

# Novelty

- What is new in this work?
  - Higher accuracy, significant speed-up, scale-up, ease to implement, generalization, wide application domain, connection among seemingly unrelated topics, ...
- What are the contributions (over prior art)?
- Make a compelling case with strong supporting evidence

# Experimental Validation

- Common data set
- Baseline experiment
- Killer data set
- Large scale experiment
- Evaluation metric
- Realize things after submission
- Friendly fire

# Compare With State of the Art

- Do your homework
- Need to know what is out there (and vice versa)
- Need to show why one's method outperforms others, and in what way?
  - speed?
  - accuracy?
  - sensitive to parameters?
  - assumption
  - easy to implement?
  - general application?

# Writing



# Writing



UCMERCED

# Writing

- Reviewing a poorly written paper
- Clear presentation
- Terse
- Careful about wording
- Make claims with strong evidence

# Writing

- Matt Welsh's blog on [scientific writing](#)
- Sharpen your mental focus
- Force you to obsess over every meticulous detail – word choice, word count, overall tone, readability of graphs (and others such as [font size](#), layout and spacing, and [page limit](#))

# Writing

- Crystalizing the ideas through the process of putting things together
- Hone the paper to a razor-sharp, articulate, polished work

# Writing

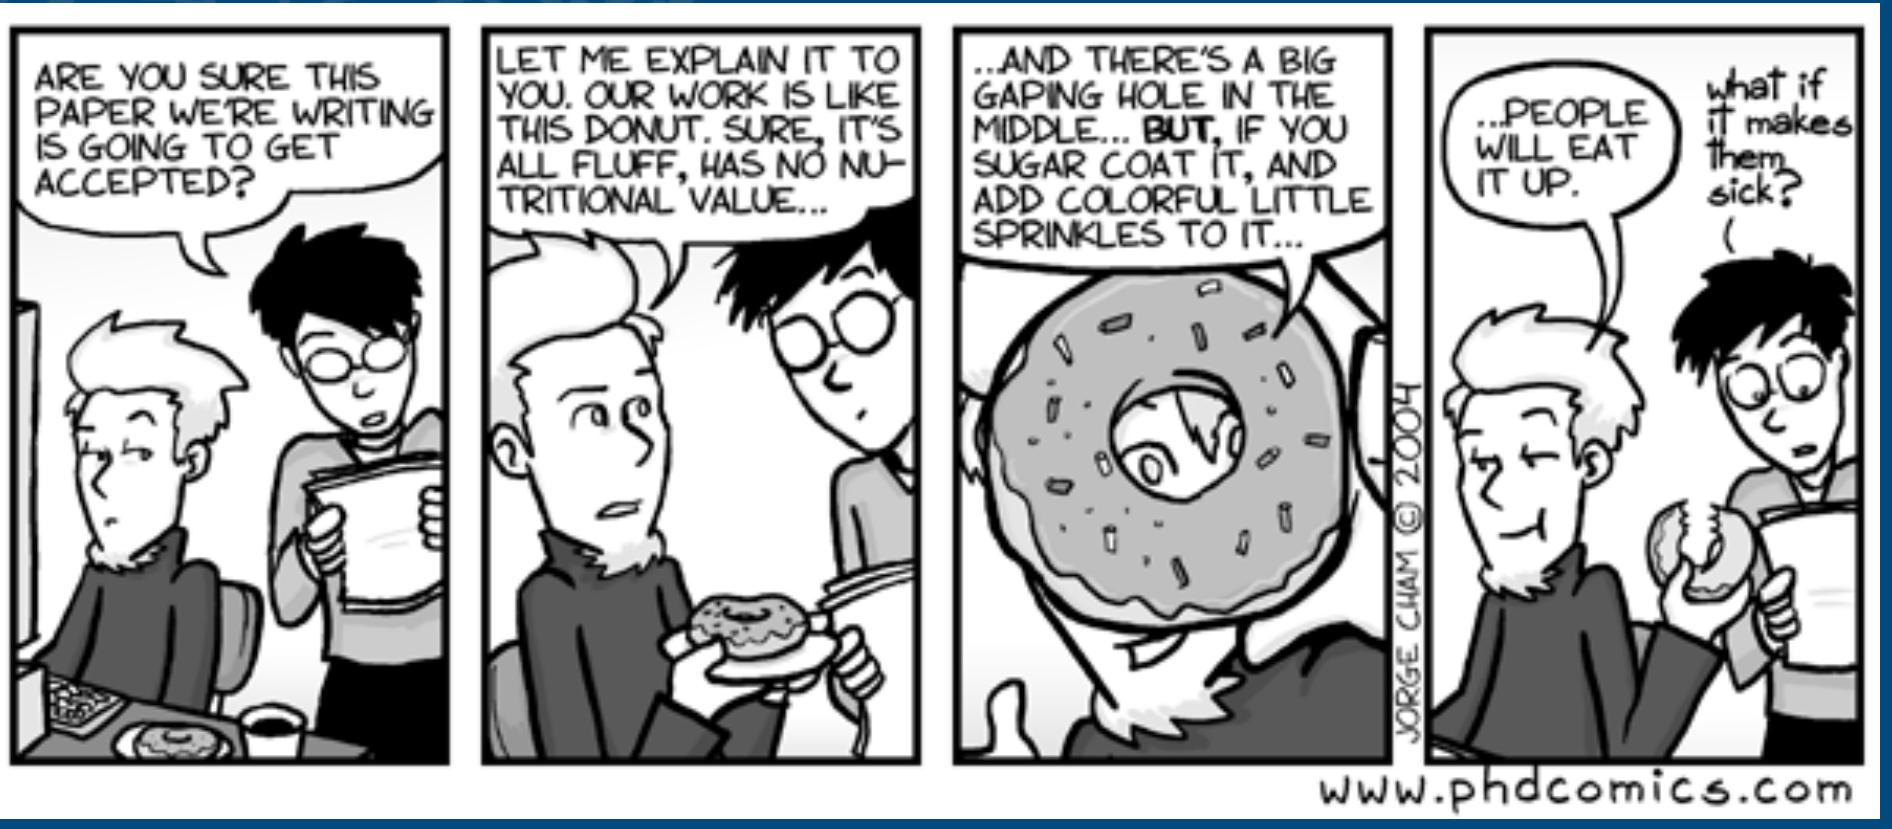
- Write the paper as *early as possible*, sometimes *before* even starting the research work
- Will discover the important things that you have not thought about
- The process of writing results in a flood of ideas

# Writing

- Even if a paper is not accepted, the process is energizing and often lead to new ideas for the next research problems
- Submitting the paper is often the start of a new line of work
- Riding on that clarity of thought would emerge post-deadline (and a much-needed break)

# Tell A Good Story

- Good ideas and convincing results
- But not too much (vs grant proposal)



# Presentation

- Good artists copy, great artists steal
- Not just sugar coating
- Not just a good spin
- Tell a convincing story with solid evidence
- Present your ideas with style
- Q&A
- Real stories

# Interesting Title

- Cool titles attract people
- Grab people's attention
- Buzz word?
- But don't be provocative

# Math Equations

- Minimal number of equations
  - No more, no less
  - Too many details simply make a paper inaccessible
- Too few equations
- Many good papers have no or few equations
  - CVPR 13 best paper
  - CVPR 05 HOG paper

# Figures

- Be clear
- Sufficient number of figures

# Theoretical or Applied?

- Computer vision is more applied, at least nowadays
- Theory vs real world
- More high impact papers are about how to get things done right

# Common Mistakes

- Typos
- Unsupported claims
- Unnecessary adjectives (superior!)
- “a”, “the”
- Inanimate objects with verbs
- Inconsistent usage of words
- Laundry list of related work (or worse copy sentences from abstracts)
- Bad references
- Laundry list of related work
- Repeated boring statements

# Get Results First than Writing?

- Conventional mode
  - Idea-> Do research -> Write paper
- “[How to write a great research paper](#)” by Simon Peyton Jones
  - Idea -> Write paper -> Do research
    - Forces us to be clear, focused
    - Crystallizes what we don't understand
    - Opens the way to dialogue with others: reality check, critique, and collaboration
- My take
  - Idea -> Write paper -> Do research -> Revise paper -> Do research -> Revise paper -> ...

# Supplementary Material

- Important
- Add more results and large figures
- Add technical details as necessary (  
don't miss important details)
- Derivation details, e.g., proof of a theorem

# Most Important Factors

- Novelty
- Significant contributions (vs. salami publishing)
- Make sure your paper is non-rejectable  
(above the bar with some error margin)

# Reviews

- Me: Here is a faster horse
- R1: You should have used my donkey
- R2: This is not a horse, it's a mule
- R3: I want a unicorn!

# Rebuttal or Response

## ADDRESSING REVIEWER COMMENTS

BAD REVIEWS ON YOUR PAPER? FOLLOW THESE GUIDELINES AND YOU MAY YET GET IT PAST THE EDITOR:

### Reviewer comment:

"The method/device/paradigm the authors propose is clearly wrong."

### How NOT to respond:

✗ "Yes, we know. We thought we could still get a paper out of it. Sorry."

### Correct response:

✓ "The reviewer raises an interesting concern. However, as the focus of this work is exploratory and not performance-based, validation was not found to be of critical importance to the contribution of the paper."

### Reviewer comment:

"The authors fail to reference the work of Smith et al., who solved the same problem 20 years ago."

### How NOT to respond:

✗ "Huh. We didn't think anybody had read that. Actually, their solution is better than ours."

### Correct response:

✓ "The reviewer raises an interesting concern. However, our work is based on completely different first principles (we use different variable names), and has a much more attractive graphical user interface."

### Reviewer comment:

"This paper is poorly written and scientifically unsound. I do not recommend it for publication."

### How NOT to respond:

✗ "You #&@\*% reviewer! I know who you are! I'm gonna get you when it's my turn to review!"

### Correct response:

✓ "The reviewer raises an interesting concern. However, we feel the reviewer did not fully comprehend the scope of the work, and misjudged the results based on incorrect assumptions."

[www.phdcomics.com](http://www.phdcomics.com)

JORGE CHAM © 2005

### Good surprise

- One CVPR paper: BR, BR, DR
- Two ECCV paper: PR, PR, BR
- One CVPR 15 paper: BR, BR, WR -> poster, poster, poster
- One CVPR 15 paper: DR, WA, BR -> Poster, Poster, WR

### Bad surprise

- Two ECCV papers: PA, PA, BR
- One CVPR 15 paper: WA, BR, BR -> Poster, Poster, WR
- One CVPR 16 paper: WR, WR, BR

# Never Know What will Happen

Masked Meta-Reviewer ID: Meta\_Reviewer\_1

Meta-Reviews:

Question

Consolidation Report

All reviewers agree that this paper has moderate novelty of using partial and spatial information for sparse representation. However, they also concern about

- unclear presentation on technical details (eg. definitions, inference algorithm, pooling methods, template updating schemes, experimental settings etc.),
- not extensive experimental comparison (needs tests on more challenging videos),
- missing justification of the assumption (complementary nature of two kinds of pooling features) and the efficacy of each term.

The authors rebuttal addresses most issues, but is not sufficient to ease the main concerns of R1 and R2. So, the AC recommends the paper to be rejected as it is.

Decision

Definitely Accept

UCMERCED

# Challenging Issues

- Large scale
  - CVPR 2011 best paper: pose estimation
  - CVPR 2013 best paper: object detection
- Unconstrained
- Real-time
  - CVPR 2001: face detector
  - CVPR 2006: scalable object recognition
- Robustness
- Recover from failure

# Interesting Stats

- Best papers and top cited papers in computer science
- Best papers = high impact?
- Oral papers are more influential?
- CVPR Longuet-Hggins prize
- ICCV Helmholtz award

# Data Set Selection

- NIPS 02 by Doudou LaLoudouana and Mambobo Bonouliqui Tarare, Lupano Tecallonou Center, Selacie, Guana
- The secret to publish a paper in machine learning conferences?
- Read the references therein carefully!

# Data Set Selection



## Data Set Selection

Donaldo LaLoudounas<sup>a\*</sup> and Mandebbo Bonouki Turare  
Lupine Technical Center  
Sofia, Bulgaria  
[dondalal@hotmail.com](mailto:dondalal@hotmail.com), [turare@yahoo.com](mailto:turare@yahoo.com)

### Abstract

We introduce the community to a new construction principle whose practical implications are very broad. It is a generalization of the well-known and widely used algorithms in data mining and to make them more appealing. We define a new notion of diversity for data sets and derive a methodology for selecting from them. The experiments show that our new approach is better than the old ones and sometimes much better than all the others. We give some experimental results, which are very promising.

### 1 Introduction

Learning is a marvellous subject. A year spent in artificial intelligence is enough to make one believe in God. Unfortunately, it has been handled only from a particular point of view. The "Catholics" have to work to prove that God exists. The others would be right to soft from such a theory: we want good bound for our algorithms. We offer such a bound here. Our bound is based on a generalization of the VC dimension. This is more principled and rigorous way that has been done before. Many researchers, especially in machine learning, have done many researches about the complexity of the function space generated by their given set of algorithms. It is strange that they ignore directly the selection problem. This is the selection problem which is what people actually do. The two problems can be seen as the dual of each other. The first problem is to estimate the error of your algorithm. And viceversa. In this article we lay down the foundations and introduce to the community of machine learning a new construction principle for data sets selection and optimization. Essentially we begin to formulate the sometimes so hazy engineering approach of data mining into a mathematical one. The new methodology gives us much better bounds for the data selected really is better than others and costs less and implement less in less time for finding such data. We leave our approach open to the closed approach.

The structure of the paper is as follows: a classical bound: section 1 presents the traditional notion of diversity for data sets and some bounds derived from this notion by defining new algorithm. Section 3 describe some experiments which, of course, are good. Section 4 concludes the article with some thoughts and future work we will do.

### 2 Bounds

Let us introduce some notations. Assume a researcher has invented an algorithm  $A^*$  and he wishes to show that his pride and joy is superior with respect to some loss function  $\ell$  to a

"Catholic" who can say a lot of things. But if we are to show that  $A^*$  is better than  $A$ , then we bounds by defining new algorithm. Section 3 describe some experiments which, of course, are good. Section 4 concludes the article with some thoughts and future work we will do.

\*Correspondence to other people, we have put all of our experience in this paper, even the bad ones, but well, we did not yet write it.

If (and only if) the paper is accepted:

Figure 1: Graphical depiction of the structural data sets minimization construction principle.

The highly successful and efficient algorithm presented in the following. Once again, assume a researcher has invented an algorithm  $A^*$  and he wishes to show that it is superior (with respect to some loss function  $\ell$ ) to a given set  $\mathcal{A}$  of algorithms  $A_1, \dots, A_k$  that other researchers have defined. Let us suppose that  $\mathcal{A}$  is a set of linear functions. Now, the task is to optimize these parameters such that the algorithm  $A$  appears much better than all others. One of the few choices is SVM, A<sub>2</sub>-SVM, A<sub>3</sub>-SVM, A<sub>4</sub>-NN, A<sub>5</sub>-CLF, etc. we choose SVM with a random distribution to run me time. This can be done with the following minimization:

$$\text{subject to: } R_{\text{emp}}^{-1}[D] \leq R_{\text{emp}}^{-1}[A_i] - \epsilon, \quad i = 1, \dots, 4$$

where  $\epsilon$  is the number of parameters (row of parameters of the vectors  $c, d$  and  $f$ ) and  $R$  is the largest possible size of the vector of any vector in the data sets.

Some people would argue that this theorem is valid only for data sets parameterized by gaussian distribution, but actually it can model a lot of real life problem. The interested reader can consult [4] for more information about this.

Now, the task is to optimize these parameters such that the algorithm  $A$  appears much better than all others. One of the few choices is SVM, A<sub>2</sub>-SVM, A<sub>3</sub>-SVM, A<sub>4</sub>-NN,

A<sub>5</sub>-CLF, etc. we choose SVM with a random distribution to run me time. This can be done with the following minimization:

which is equivalent to

$$\min_{w, c, d, f} \|w\|^2$$

i.e., slightly altered source's.

It is now clear why we said before, we repeat the shortest introduction. The problem with our previous lies is that you don't know where to go to create it. This is the same for other metrics by the way. Consider for instance the notation. It is not at all clear where it comes from, or how it is related to the rest. The same also holds for terms mostly when the authors are new in the field. Note that we also are in the field but we are not the first. We are not the second either. We are not the third or the fourth or even the last. Anyways, the question we are discussing right now is to know when and how to stop. The answer is not in the table, it is not in the table. We are not the best in the world of course, we will not reiterate the introduction. This may then sound weird that we, as engineers, put a note in the end of the paper. We do not know the fundamentals level of each of the researchers, but we do know that some of them are not very good.

We will now explain what we said before, we repeat the shortest introduction.

The problem with our previous lies is that you don't know where to go to create it. This is

the same for other metrics by the way and come back to our contribution. Central to our research is the idea to impose the presentation of algorithms in literature and to make them more appealing. We also try to make them more accessible for non-experts in the methodology. The experiments showed that even if not so good algorithm, you can show that they are significantly better than all the others. The message is strong and may not be heard by a few. We hope that our research will help to bring the message to a wider range of researchers to dig out their old failed algorithms and turn them into successful ones.

given fixed set of algorithms  $A_1, \dots, A_k$  that other researchers have made. For this purpose, the researcher selects some data sets using what is called an empirical data set minimization method. The data sets can be found in the Internet and are often given  $D_1, \dots, D_k$  and find it data set  $D^*$  in  $D_1, \dots, D_k$  so that:

$$R_{\text{emp}}^{-1}[D^*] < \min_{i=1, \dots, k} R_{\text{emp}}^{-1}[A_i]$$

Now that this problem is ill-posed. A natural continuation would be to find more than one data set in which your algorithm performs well but this is a difficult problem that has not been solved by us so far. Current efforts in solving this problem have focused on random sampling and some specific search methods.

We have the following theorem:

**Theorem 1.** For any  $\epsilon > 0$  of running rate, that ensures that the prior of algorithms is enclosed with a fixed interval  $I$  which could be  $[0, 1)$  or  $[0, \infty)$ , there is probability  $1 - \eta$  a sampling on the  $D^*$  such that

$$V[D] \in D, \quad R_{\text{emp}}^{-1}[D] \leq R_{\text{emp}}^{-1}[D^*] + O(\sqrt{\frac{\log(1/\eta)}{m}})$$

where  $\Theta(D)$  is the property of the set of training data defined as:

$$\Theta(D) = \max_{\text{w.r.t. } A_1, \dots, A_k} \text{Algorithm } x \cdot \ell(A_1, \dots, A_k), \quad \forall x \in [0, 1]^{\sum_{i=1}^k |A_i| - 1}, \quad (1)$$

We are proud now to supply the following elegant proof.

**Proof.** We consider  $n$  as the number of points in the training set, we see that introducing a global algorithm  $A$ :

$$P\left(\sup_{A \in \Theta(D)} R_{\text{emp}}^{-1}[D] - R_{\text{emp}}^{-1}[D^*] > \epsilon\right) \leq P_{\text{exp}}\left(\sup_{A \in \Theta(D)} R_{\text{emp}}^{-1}[D] - R_{\text{emp}}^{-1}[D^*] > \epsilon\right)$$

which is trivially insensitive to permutations so that we can condition the algorithms  $A$  on  $D$ . So, we also have the right to perform with non-permuting permutations as we have done here. We also use the VC dimension and the VC function  $\Phi(D)$  so that we work only with the values of  $(c_1, c_2)$ . After some more states which we admit for brevity that the function  $\ell(A)$  is bounded by the function  $\ell$  for all the functions in the class  $A$  and the random variables which can be controlled using the Bennett-Verma's inequality. Thus  $\ell(A)$  is bounded by the function  $\ell$  for all the functions in the class  $A$ . Each of these variables does not give you a good control of their expectation. But this can be overcome by using the mean value theorem. The mean value theorem gives the value of  $\ell$  as the midpoint of  $I$ . As we can see, the expectation gives the value of  $\ell$  as the midpoint of  $I$ . So, we can use the mean value theorem to find the value of  $\ell$  as the midpoint of  $I$ . We call this trick, the replica trick. Now the replica trick has made us need more knowledge in the area of physics and it is not for everyone. So we leave it to all possible algorithms, the new algorithms have been tested many times with negligible variations with respect to the original one.

The theorem we just proved should be considered as the dual of the theorem of Vapnik and Chervonenkis. And this should be the case because it is not the dual of  $\ell$ . And we believe that the main motivation for your experiments and its applications are the same as ours (so we can apply the same to the other). But the reason for this is that Vapnik and Chervonenkis are one step far from us because we can fully compute the probability over the data sets just by looking at the LCI criterion. We can say that the LCI criterion is the most important criterion for data sets. Anyways, we insist that our approach appeals a lot of common properties with the classical VC dimension. The reader can see that the two theorems are very similar. But let us say differently we have our own "no free lunch" theorem but we try to forget it. Here is our proof of the theorem:

\*Actually, a small set so that that she does not have to do too many experiments. In fact, coding nice paper into a vector space, we found a big cluster where all the points were close to each other as well as they were far away. We can use the same approach throughout and as statistics. We can't assume 50 times but kept coming up with the same single big cluster.

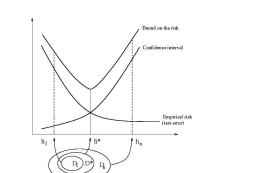


Figure 1: Graphical depiction of the structural data sets minimization construction principle.

| Algorithm  | Worst-case      | Worst-case     |
|------------|-----------------|----------------|
| Poly-2 SVM | 0.5 (0.1)       | 0.5 (0.2)      |
| Poly-4 SVM | 0.3 (0.1)       | 0.2 (0.1)      |
| TKL-2 SVM  | 0.9 (0.9)       | 0.8 (0.6)      |
| TKL-4 SVM  | 0.9 (0.9)       | 0.8 (0.6)      |
| Recto      | 0.001 (0.00009) | 0.002 (0.0001) |

Figure 2: Results for  $w = (0.00001, 0.0000145839, 34, 31, 4120, 2.5, 2.3, 5.6, 4.7, 8.9, -1, -2, -3, -3)$ , w corresponds to the parameters of the generated data set. Unfortunately, we have difficulties interpreting on.

$$\text{subject to: } R_{\text{emp}}^{-1}[D] \leq R_{\text{emp}}^{-1}[A_i] - \epsilon, \quad i = 1, \dots, 4$$

If you wish to find a dataset where your algorithm does not achieve 0% test error you can easily generate this algorithm to the linearly inseparable case by introducing vector variables  $\zeta$ .

The classes to SVM is striking. Note that we are maximizing the margin between the algorithms to ensure a paper is accepted. The relation with a paper is open and has no boundary. Now we are talking about a lot of work in statistical test selection so that even with small margins you can have a strong result.

### 5 Experiments

We proceed to prove our methodology by showing we can always find a good solution even for bad algorithm. So we proceed with the example given above. We must admit that it has already been done by other authors in different ways, but we present it here.

In table 1, we present the results we get with the best value for  $w$ .

Note that we have the same behavior as in figure 1. We can see that the poly-2 SVM looks like it would have a similar behavior and we also want to worsen the results of the poly-2 SVM although it would have been nice. Poly-2 SVM prefers to work on large number of data sets and this is the reason why we can not use it as an algorithm for a paper. So on the other hand, it is quite clear in the table that Recto is much better than all the others. We can also see that the TKL-2 SVM and TKL-4 SVM have some good common properties with Recto. Thus our algorithm was able to discriminate between 2 distances in the space and we can expect to work very well for the neural networks.

### 6 Conclusion

We will now estimate what we said before, we repeat the shortest introduction. The problem with our previous lies is that you don't know where to go to create it. This is the same for other metrics by the way. Consider for instance the notation.

We have the following theorem<sup>2</sup>: The generalization error of two datasets for all algorithms is the same:

$$E[R_{\text{emp}}^{-1}[D]] = E[R_{\text{emp}}^{-1}[D^*]]$$

so that there is no better dataset than the one you already have.

The consequences of the theorem are very hard in the sense that if someone had a break, then you are not able to continue with the algorithm with which someone had a break (of course, in this remark we retain some liberty). This leads to the natural consequence that we can't continue our work without starting from scratch. In our companion paper we prove that the set of datasets restricted to all Bayesian algorithms has no such dimension. The same is true for the set of datasets restricted to all kernel algorithms if we have the following conditions. The same is true for the set of datasets restricted to all bayesian algorithms (see the following theorem).

Theorem 2 (No Free Lunch!) The generalization error of two datasets for all algorithms is the same:

$$E[R_{\text{emp}}^{-1}[D]] = E[R_{\text{emp}}^{-1}[D^*]]$$

so that there is no better dataset than the one you already have.

The consequences of the theorem are very hard in the sense that if someone had a break, then you are not able to continue with the algorithm with which someone had a break (of course, in this remark we retain some liberty).

This leads to the natural consequence that we can't continue our work without starting from scratch. In our companion paper we prove that the set of datasets restricted to all Bayesian algorithms has no such dimension. The same is true for the set of datasets restricted to all kernel algorithms if we have the following conditions.

Theorem 3 (Structural Data Sets Minimization) In order to appreciate the following section, we ask the reader for a little patience and to conceive that our data sets  $D_1, \dots, D_n$  (e.g. USPS, NYST, REUTERS 1, REUTERS 2, ...) which are roughly included in the datasets  $D_1, \dots, D_n$  (e.g. USPS, NYST, REUTERS 1, REUTERS 2, ...)

Then we would like to apply the theorem of the previous section and choose the best datasets for our algorithm (e.g. the one that we can use for presenting our paper). Using the name based we found that with probability  $1 - \sum_{i=1}^n \alpha_i$  for all  $D \in D_1, \dots, D_n$ :

$$R_{\text{emp}}^{-1}[D] \leq \min_{D \in D_1, \dots, D_n} R_{\text{emp}}^{-1}(D) + O\left(\sqrt{\frac{\Phi(D)}{m} \log(1/\alpha)}\right)$$

So we can use the  $i^{th}$  (and that  $R_{\text{emp}}[i]$ ) is minimized over  $i$  and choose then the best data set in  $D_1$ . The consistency of this procedure can be proved by a theorem.

This section means that to find the best data set for your algorithm, you can consider many datasets (not necessarily your data set) and then choose the one that has the lowest expected error (which is extremely often called one error in some papers) but also to trust this data set.

Here is another picture (Figure 1).

### 4 Algorithm Ordering Machine

We introduce a general method for choosing between random elements. However, we do not really know the capacity of these real world datasets, only approximations can be measured through empirical observations, that most probably is not good enough to do well so we capacity must be quite large. For this reason we use the toy problem. It turns out that we can use an efficient algorithm for searching for the best algorithm.

It is a well-known algorithm, the Replicon. We hope that the reader will learn from this breakthrough and applies our methodology in their future research or they will get fully behind our algorithms for jet selection problems.<sup>3</sup>

### References

- [1] R. Schapire, Y. Freund, P. Bartlett, and W.S. Lee, *Balancing for margins*. *Annals of Statistics*, 2009.
- [2] N. Cristianini and J. Shawe-Taylor, *Data Set Selection for Demeans*. Egham University Press, 1996.
- [3] P. Bartlett, *The Sample Complexity of Pattern Classification with Neural Networks: the Number of Classes is More Important than the Number of Nodes*. *Bulletin of the American Mathematical Society*, 1998.
- [4] D. Moshkovitz, *I did it my way*. *Sifir published*, 2002.
- [5] D. Slepcev, *Concentration of Measure on Product Spaces*. *Probability Theory and Related Fields*, 2000.
- [6] J. Friedman, T. Hastie, and R. Tibshirani, *Pathwise coordinate optimization*. *Technometrics*, 2008.
- [7] D.M. Witten, *Statistical Topic Selection in Linear Discrimination*. *Oxford University Press*, 1995.
- [8] D.H. Wolpert, *The lack of a priori distinction between learning algorithms revisited*. *Neural Computation*, 1996.
- [9] N.S. Vapnik, and V. Chervonenkis, *On the Uniform Convergence of Frequency of Success to its Expectation*. *IEEE Transactions on Information Theory*, 1963.
- [10] J. Kittler, J. H. Wolf, and J. M. Nasrabadi, *The VC Way: Incremental Secrets from the Wizards of Venice Capital*. BioWorld Publishing (now out of print), 2001.

<sup>2</sup>Note that for us, since people conjecture it should be in some unknown place, others yet conjecture it doesn't exist at all. We will refer to this as the Dijet and note that at least it is very hard to find, at least if the  $\ell$ 's are very small.

<sup>3</sup>Note how many people try to implement random forests, it is a very good idea, requires intelligence, vigilance, dedication and courage. If we don't practice these traits of thought, we cannot hope to solve the toy problem. Note that this is a very good solution for jet selection problems.

<sup>4</sup>Note that this is a very good solution for jet selection problems that face us – and we risk becoming a series of readers, as for guide by the next deadlines who comes along.

# Data Set Selection

## Data Set Selection

**Doudou LaLoudouana\* and Mambobo Bonouliqui Tarare**

Lupano Tecallonou Center  
Selacie, GUANA

doudoula3@hotmail.com, fuzzybear@yahoo.com

### Abstract

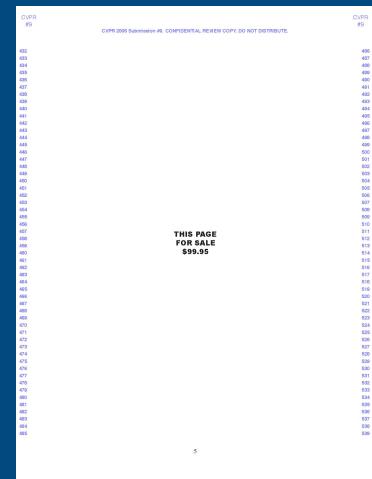
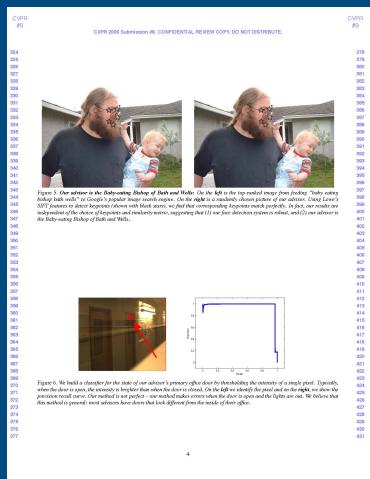
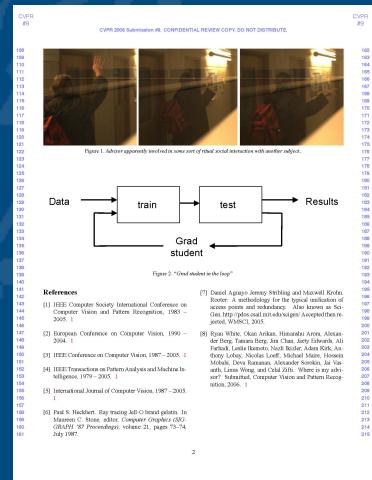
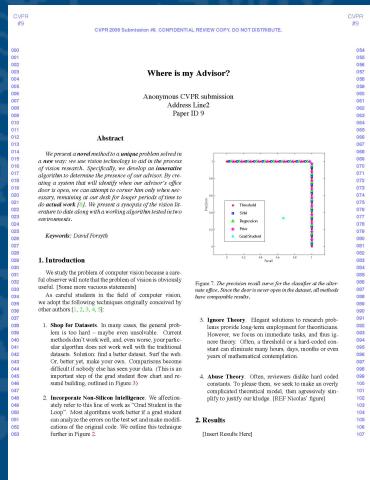
We introduce the community to a new construction principle whose practical implications are very broad. Central to this research is the idea to improve the presentation of algorithms in the literature and to make them more appealing. We define a new notion of capacity for data sets and derive a methodology for selecting from them. The experiments show that even for not so good algorithms, you can show that they are significantly better than all the others. We give some experimental results, which are very promising.

# Data Set Selection

## References

- [1] R. Shapire, Y. Freund, P. Bartlett, and W.S. Lee. *Bushing the margin: A new explanation for the effectiveness of U.S voting methods.* *The Annals of Statistics*, 1998.
- [2] N. Cristianini and J. Shawe-Taylor. *Data Set Selection for Dummies*. Egham University Press, 2005.
- [3] P. Bartlett. *The Sample Complexity of Pattern Classification with Neural Networks: the Number of Citations is More Important than the Number of Readers.* *Biowulf Transactions* , 1998.
- [4] D. MacKay. *I did it my way.* Self published, 2002.
- [5] G. Fung and O. L. Mangasarian. Data Selection for Support Vector Machine Classifiers. Proceedings of KDD'2000, 2000.
- [6] J. Friedman, T. Hastie, and R. Tibshirani. *Data sets for additive models: a statistical view of bragging and boasting.* Technical report, Technical Report, Department of Statistics, Stanford University., 1998.
- [7] C.M. Bishop. *Neural Networks for Pattern Recognition*, Oxford University Press, 1995
- [8] D.H. Wolpert, *The lack of a priori distinctions between learning algorithm research*, Neutral Computation, 1996.
- [9] V.N. Vapnik and A.Y Chervonenkis and S. Barnhill, *The VC Way: Investment Secrets from the Wizards of Venture Capital*, Biowulf Publishing (now out of print) , 2001.

# Where Is My Advisor?

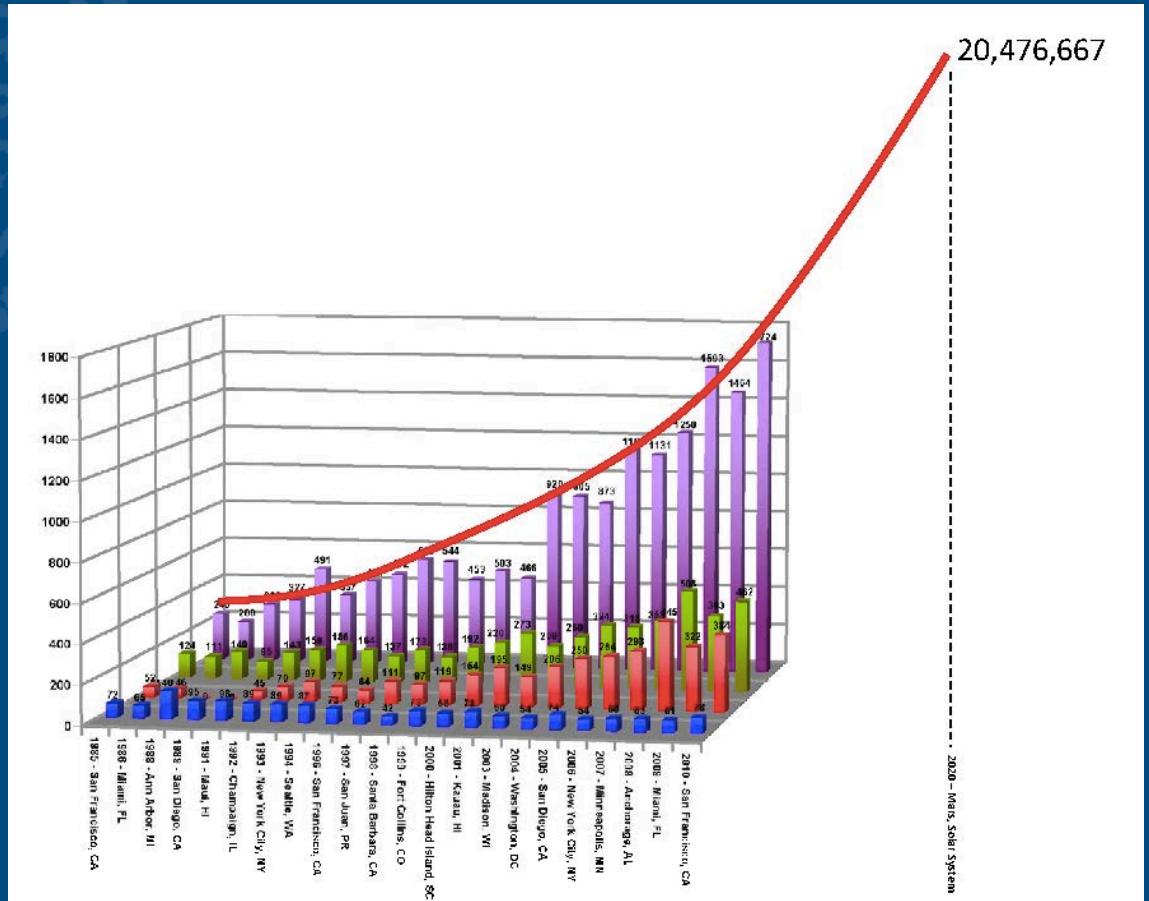
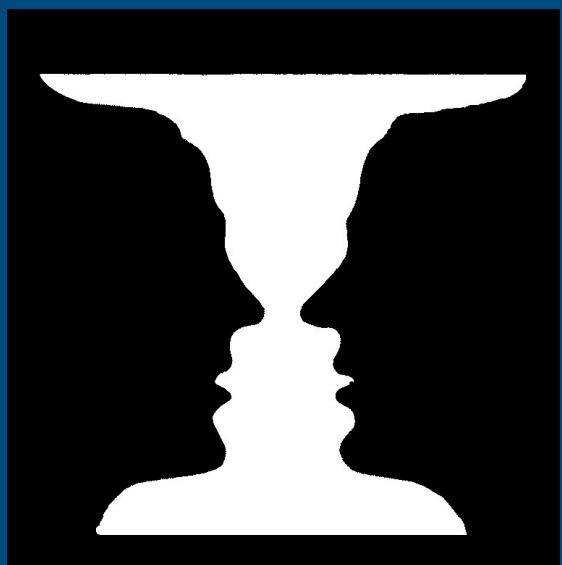


UCMERCED

# Ask Someone to Proofread

- Certainly your advisor
- Polish your work
- My story

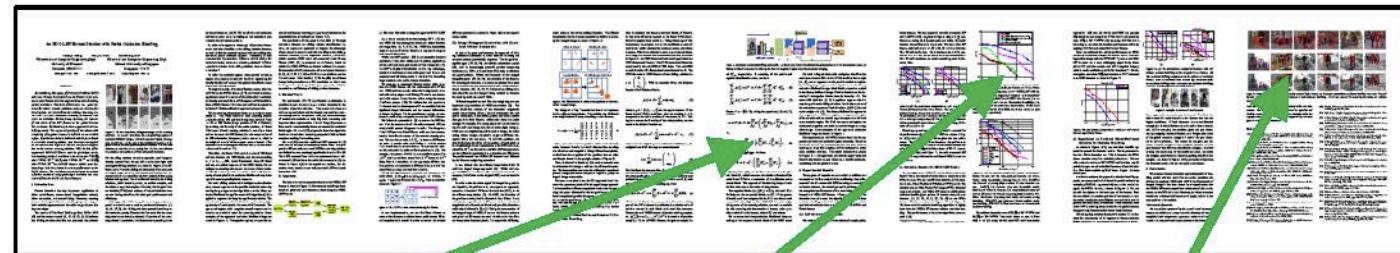
# Paper Gestalt



**Figure 1. Paper submission trends.** The number of submitted papers to CVPR, and other top tier computer vision conferences, is growing at an alarming rate. In this paper we propose an automated method of rejected sub-par papers, thereby reducing the burden on reviewers.

# Paper Gestalt

- [CVPR 10 by Carven von Bearnensquash, Department of Computer Science, University of Phoenix](#)
- Main Point: Get your paper looking pretty with right mix of equations, tables and figures

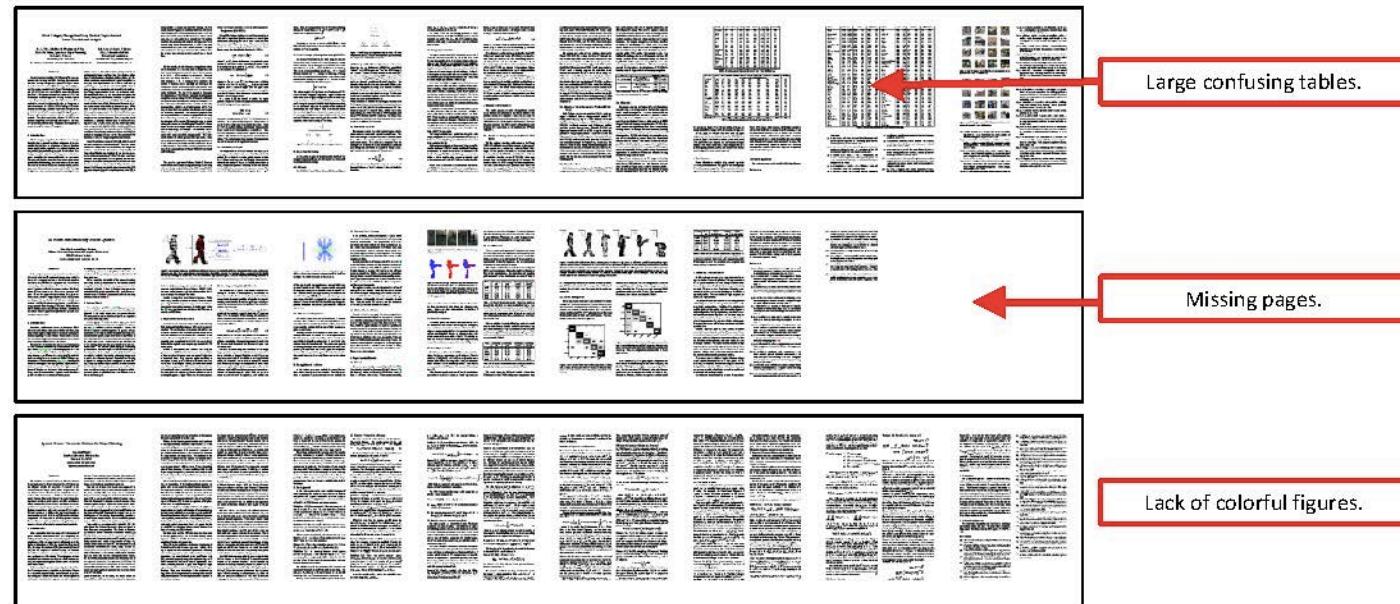


**Math:** Sophisticated mathematical expressions make a paper look technical and make the authors appear knowledgeable and "smart".

**Plots:** ROC, PR, and other performance plots convey a sense of thoroughness. Standard deviation bars are particularly pleasing to a scientific eye.

**Figures/Screenshots:** Illustrative figures that express complex algorithms in terms of 3<sup>rd</sup> grade visuals are always a must. Screenshots of anecdotal results are also very effective.

Figure 6. Characteristics of a "Good" paper.



Large confusing tables.

Missing pages.

Lack of colorful figures.

Figure 7. Characteristics of a "Bad" paper.

# Tools

- Google scholar h-index
- Software: publish or perish
- DBLP
- Mathematics genealogy
- Disclaimer:
  - h index = significance?
  - # of citation = significance?

# Basic Rules

- Use LaTeX
- Read authors' guideline
- Read reviewers' guideline
- Print out your paper – what you see may NOT be what you get
- Submit paper right before deadline
  - Risky
  - Exhausting
  - Murphy's law
- Do not count on extension

# Lessons

- Several influential papers have been rejected once or twice
- Some best papers make little impact
- Never give up in the process

# Karma?

## WHAT YOU THINK OF YOUR PROFESSOR vs. TIME



JORGE CHAM © 2014

[WWW.PHDCOMICS.COM](http://WWW.PHDCOMICS.COM)

UCMERCED

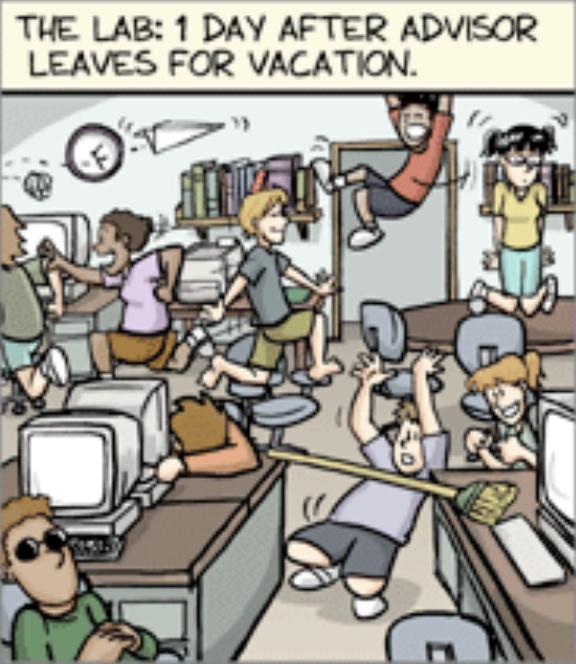
# Your Advisor and You

- Suggesting a research topic
- When your advisor presents your work
- When you explain your work
- Demos
- Good results

# Start Working Early!

- Write, write, write...
- Ask others for comments

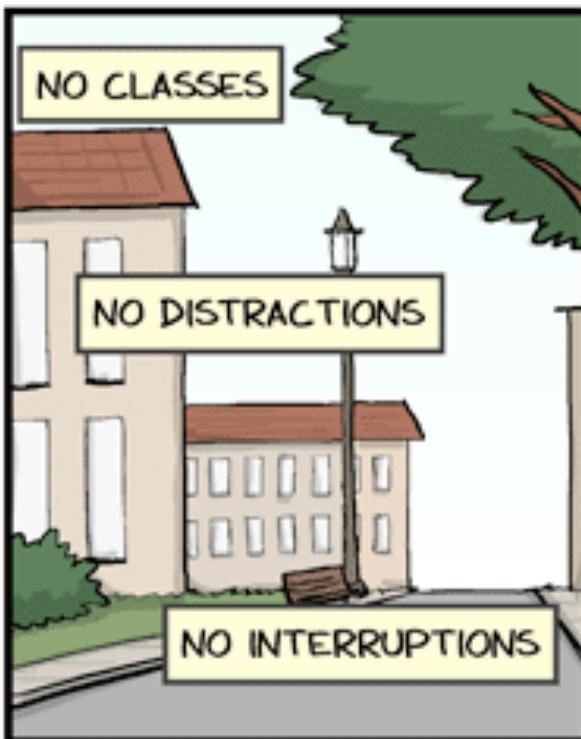
## SUMMER DAYS...



JORGE CHAM © 2004

[www.phdcomics.com](http://www.phdcomics.com)

# Work Hard in the Summer



UCMERCED

# Quotes from Steve Jobs

- “ I'm convinced that about half of what separates successful entrepreneurs from the non-successful ones is pure perseverance. ”
- “ Creativity is just connecting things. When you ask creative people how they did something, they feel a little guilty because they didn't really do it, they just saw something. It seemed obvious to them after a while. ”